

## What are BLOBS?

BLOB = Binary Large Object

A BLOB is raw binary data, for example:

- Images (JPEG, PNG)
- Videos (MP4)
- Audio (MP3)
- PDFs
- ZIP files

In databases, BLOBS are stored as:  
BLOB, BYTEA, VARBINARY, etc.

```
CREATE TABLE users (
    id INT,
    profile_pic BLOB
);
```

At first glance, this looks convenient:

"User data + image in one place"

But this breaks down very fast at scale.

## Why Storing BLOBS in a Database Is a Bad Idea ?

### 1. Databases Are Optimized for Structured Data

Databases are built for:

- Rows & columns
- Indexes
- Queries
- Joins
- Transactions

They are not optimized for large binary payloads.

When you store BLOBS:

- Indexes become useless
- Pages get bloated
- Cache efficiency drops

You're using a Ferrari as a truck.

### 2. Performance Problems

Let's say:

Profile image = 5 MB

User table row = 1 KB

Now every time you do:

```
SELECT * FROM users WHERE id = 123;
```

The DB:

- Reads the entire row
- Pulls the BLOB into memory
- Sends it over the network
- Even if you don't need the image.

Result:

- Higher latency
- More memory usage
- Slower queries

### 3. Backup & Restore Become Nightmares

Databases need:

- Regular backups
- Replication
- Point-in-time recovery

With BLOBS:

- Backup size explodes
- Restore takes hours
- Replication lag increases

### 4. Scaling Becomes Extremely Expensive

To scale a DB:

- Vertical scaling (bigger machine)
- Sharding (very complex)

Binary data:  
Increases storage cost  
Increases I/O  
Increases replication traffic

DBs are the expensive storage in your system.

Using them for images/videos is burning money.

### 5. Databases Are Stateful & Coupled

Databases:

- Are tightly coupled to applications
- Cannot be easily accessed directly by clients
- Don't integrate well with CDNs

You don't do:



That alone disqualifies DBs for media delivery.

The Correct Pattern (Industry Standard)

```
Image → Object Storage  
Metadata → Database
```

## Why We Need Object Storage ?

We want storage that is:

Cheap

Scalable

Durable

Accessible over HTTP

CDN-friendly

Not tied to DB constraints

This is exactly why Object Storage exists.

## What Is Object Storage?

Object Storage stores data as objects, where each object contains:

Data (binary content)

Metadata (key-value info)

Object ID (key)

Objects live inside buckets and are accessed using HTTP APIs.

Example:

bucket: user-uploads

key: profile-pics/aniket.png

## Examples of Object Storage

Amazon S3

Google Cloud Storage

Azure Blob Storage

## Upload Flow (PUT Object)

PUT /user-uploads/profile-pics/aniket.png


Client

(Browser / App)

API Layer  
Auth  
Validation

Metadata Service

create entry  
(PENDING)



Step 1: Request Comes In

API Layer checks:

Auth (IAM / token)

Bucket exists?

Permission?

Size limits?

Step 2: Metadata Entry Created

Why PENDING?

Upload not finished yet

Metadata DB:

key	size	status	locations
img1	5MB	PENDING	NULL

### Step 3: Object Is Chunked

Object (20MB)



Chunks allow:

Parallel upload

Retry on failure

Large object support

### Step 4: Chunks Distributed to Nodes


### Step 5: Durability Applied (Replication)

### Step 6: Metadata Updated (COMMITTED)

Metadata DB:

key	size	status	locations
img1	5MB	READY	NULL

Only now client gets 200 OK.



### Step 1: Lookup Metadata

key → [Node A, Node B, Node C, Node D]

### Step 2: Fetch Chunks (Parallel)

Node A → c1  
Node B → c2  
Node C → c3  
Node D → c4

Data Nodes


### Step 3: Reassemble Object

c1 + c2 + c3 + c4 = Full Object

### Step 4: Stream to Client

### Delete Flow

DELETE object



Actual data deletion:

Happens asynchronously

Garbage collector removes chunks later

## Metadata Is the Brain

Metadata Service  
└─ Object key  
└─ Size  
└─ Version  
└─ ACL  
└─ Checksums  
└─ Data locations

If metadata is down:  
Storage is down (even if data exists)

## Versioning (Overwrite Case)

PUT img.png (v1)  
PUT img.png (v2)  
Metadata:  
img.png → v2 (active)  
v1 (older)  
Delete:  
DELETE img.png  
→ delete marker added

## Multipart Upload (Large Files)

### Initiate Upload

```
|  
Upload Part 1      Part1 --> Node A  
Upload Part 2      Part2 --> Node B  
Upload Part 3      Part3 --> Node C  
|  
Complete Upload
```

Benefits:

Resume

Parallelism

Reliability

## Core Characteristics

### 1. Immutability

Objects cannot be modified

You overwrite or delete + re-upload

### 3. Globally Unique Object Key

Bucket + object key

Example:

s3://user-uploads/profile-pics/aniket.png

### 2. Flat Namespace

No real folders  
/photos/2025/img.png is just a string key

Folders are:

UI illusion  
Prefix-based filtering

## Durability (Why Your Data Never Gets Lost)

Object storage achieves insane durability using:

-> Replication

Multiple full copies

Simple but costly

-> Erasure Coding (Preferred)

Data split into  $k + m$  blocks

Lose some blocks → still recover

This is how services like Amazon S3 achieve 99.99999999% durability (11 nines).

## Security

Encryption

At rest (AES-256)

In transit (TLS)

Access Control

Bucket policies

Object ACLs

IAM roles

Pre-Signed URLs

Temporary access

Used for uploads/downloads without exposing credentials

## Performance Characteristics

Not for:

Low-latency random writes

Frequent updates

Databases

Good for:

Large objects

Streaming

Read-heavy workloads

CDN integration

## Common Use Cases

Media hosting (video/audio/images)

Backups & snapshots

Logs & analytics

Data lakes

ML training datasets

## Tips for Interviews and real-world systems

1. Never store large files in a database



Instead, we split responsibilities:

Object Storage

Stores:

Images

Videos

PDFs

Large binary files

Example:

Post image

Video file

Backup file

Database

Stores metadata only:

Post

- id

- creatorId

- text


- linkToObjectStore

That linkToObjectStore is usually:

S3 object key

Or a URL

## 2. Direct upload to Object Storage using pre-signed URLs



1. Client asks server:  
"I want to upload a file"

2. Server:  
Authenticates user  
Generates a pre-signed URL from object storage

3. Server sends URL back to client

4. Client uploads directly to object storage

This is how:


Instagram  
Netflix uploads  
Google Drive  
work.

In real systems

With Amazon S3:

Server creates a pre-signed URL  
Browser uploads directly to S3  
Server never touches the file bytes

Bad approach:



Problems:

Server bandwidth bottleneck  
High memory usage  
Poor scalability

## 3. Large files are uploaded in chunks (multipart upload)

Why chunked upload is required

Large files (like 5–10 GB):

Cannot be uploaded reliably in one go  
Network failures are common

So object storage:

Accepts file in multiple parts  
Stores parts independently  
Combines them at the end