


## Why WebSockets Came (What was broken in HTTP)

Original HTTP Model (Request → Response)

HTTP was designed for documents, not real-time systems.

- Basic rule of HTTP:
- > Client sends a request
  - > Server sends a response
  - > Connection closes (HTTP/1.1: mostly)



That's it.

Problem:

What if the server has new data, but the client didn't ask?

-> Server cannot push data.

## Real-world Use Cases That Broke HTTP

Modern apps need real-time:

Chat apps (WhatsApp, Slack)


Stock price updates

Notifications

Multiplayer games

Live dashboards

Online collaboration (Google Docs)



HTTP alone cannot do this efficiently.

## Hacks Before WebSockets (How we tried to fix HTTP)

### 1. Polling (The worst but simplest)

Client asks server every X seconds


"Any update?"

Problems:

- > Huge number of useless requests
- > High latency (updates delayed until next poll)
- > Wasted CPU, bandwidth, money

Example:

10M users × poll every 5 sec = 120M requests/minute



## 2. Long Polling (Slightly better)


- > Client sends request
- > Server waits until data is available
- > Server responds
- > Client immediately sends another request

Problems:

Still HTTP overhead (headers again & again)

Connection constantly created & destroyed

Hard to scale



## 3. Server-Sent Events (SSE)

- > Client opens one HTTP connection


- > Server pushes updates

Limitations:

-> One-way only (server → client)

-> No binary data

-> Not suitable for chat/games



HTTP was never meant for real-time

Why WebSockets Were Invented ?

Goal of WebSockets

Create a protocol that:

- > Keeps one persistent connection
- > Allows full-duplex communication
- > Has low overhead
- > Works well with browsers
- > Plays nicely with existing HTTP infrastructure


WebSockets were standardized

### What Exactly Is WebSocket?

A persistent, full-duplex communication protocol over a single TCP connection

Key Properties:

- > One connection
- > Client ↔ Server both can send anytime
- > Very low latency
- > Minimal overhead



## How WebSocket Connection Is Established (Handshake)

### Step 0 – TCP connection (happens silently)

Before anything else:

-> Client opens TCP connection to server  
-> Port 80 (ws) or 443 (wss)

-> If wss → TLS handshake happens


This is normal TCP/HTTPS, nothing special yet.

### Step 1 – Client sends HTTP request (Upgrade request)

Only the CLIENT can start the upgrade

Client sends a normal HTTP request with special headers:

```
GET /chat HTTP/1.1
Host: example.com
Upgrade: websocket
Connection: Upgrade
Sec-WebSocket-Key: random_key
Sec-WebSocket-Version: 13
```



What each header means

Upgrade: websocket  
-> "I want to switch protocol"

Connection: Upgrade  
-> "This connection should be upgraded"


Sec-WebSocket-Key  
-> Random value (security check)

Sec-WebSocket-Version: 13  
-> WebSocket protocol version

Important:

This is still pure HTTP

No WebSocket frames yet



### Step 2 – Server validates the request

Server checks:

- > Is Upgrade: websocket present?
- > Is version 13 supported?
- > Is this endpoint allowed to upgrade?
- > If any check fails → server replies with normal HTTP error

### Step 3 – Server sends HTTP 101 response (Upgrade accepted)


### Step 4 – Protocol switches (this is the magic moment)

After 101 response:

- > No more HTTP
- > WebSocket protocol starts

From this point:

No HTTP headers  
No request/response model  
Only WebSocket frames



## Step 5 — WebSocket communication begins


Now both sides can send:

- Text frames
- Binary frames
- Ping / Pong
- Close frames

And:

Server can send without waiting  
Client can send anytime

-> This is full-duplex



## How a WebSocket Connection Closes — Step by Step

WebSocket closes via a small, graceful handshake using CLOSE frames — not by just killing TCP.

### Step 0 — Connection is active

TCP connection is open  
WebSocket frames are flowing (text/binary)  
Ping/Pong may be happening

### Step 1 — One side decides to close

Either side can start:

Client (browser refresh, tab close)  
Server (shutdown, idle timeout, deploy)

-> There is no "request" or "response" here — only frames.

### Step 2 — Initiator sends a CLOSE frame

The initiator sends a WebSocket CLOSE frame.

CLOSE frame contains:

Close code (optional but recommended)  
Reason string (optional)

Example:

Opcode: CLOSE  
Code: 1000  
Reason: Normal closure

At this moment:

Initiator will NOT send any more data frames

Only allowed frames after this:  
→ Close / Pong (protocol rule)

### Step 3 — Receiver gets CLOSE frame

The other side receives the CLOSE frame and must:

-> Stop sending normal data

-> Send its own CLOSE frame back  
(echoing or choosing its own code)

This is mandatory for a graceful shutdown.

### Step 4 — TCP connection is closed

After both sides have exchanged CLOSE frames:

-> TCP connection is closed

-> File descriptors released

-> Memory cleaned up

--> WebSocket is fully closed

--> No resource leak

Client — CLOSE —> Server  
Client <- CLOSE —> Server  
TCP connection closed


## Scaling WebSockets

Problem:

WebSocket is stateful  
Connection tied to one server  
Load balancer must use sticky sessions

Scaling Solutions:

- > Sticky load balancing
- > Shared pub/sub (Redis, Kafka, etc.)
- > Connection registries
- > Horizontal scaling with event fan-out



## Failure Handling in WebSockets

Network drops → connection lost

Server crash → clients disconnect

Common strategies:

- > Heartbeats (Ping/Pong)
- > Auto reconnect on client
- > Resume sessions
- > Message acknowledgements (app-level)
- > WebSocket itself does not guarantee delivery

## When You SHOULD Use WebSockets

- > Chat systems
- > Live notifications
- > Real-time dashboards
- > Online games
- > Collaborative editors
- > Live trading apps

## When You SHOULD NOT Use WebSockets

- > Simple CRUD APIs
- > Rare updates
- > Stateless REST services
- > If SSE is enough (one-way)

DEMO