RabbitMQ

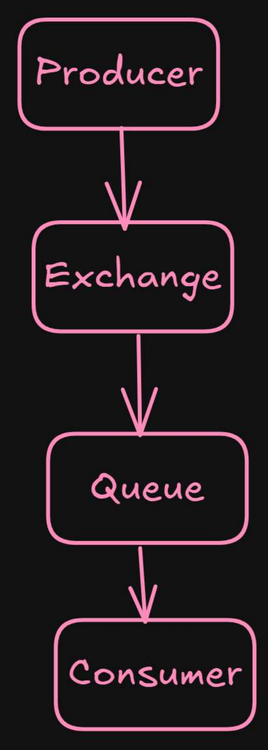

RabbitMQ is an open-source message broker (aka message queue).
It helps different parts of a system communicate asynchronously and decouple producers and consumers.

Think of it as a post office:

You (Producer) put a letter (Message) into the postbox (Queue)

RabbitMQ (Post Office) ensures it's delivered

Someone else (Consumer) picks it up when ready

Main job: Receive messages, route them, store them temporarily, and deliver them to consumers reliably.

Why do we need RabbitMQ?

Let's take an example:
You place an order on Amazon.

The app needs to:

Confirm order via email
Update inventory
Charge payment
Notify delivery service

Doing all this synchronously (in one request) would be slow and error-prone.

Instead, the Order Service publishes events ("Order Placed") to RabbitMQ.
Other services (Email, Inventory, Payment, Delivery) consume the message asynchronously.

If any service is down, RabbitMQ keeps the message safe.

Benefits

Decoupling (services don't depend directly)

Scalability (add more consumers easily)

Reliability (messages persist even if services crash)

Asynchronous processing

Load balancing

Retry and Dead-lettering support

Core Components

1. Producer

Sends (publishes) messages to an Exchange.

It never sends messages directly to a queue.
Instead, it always sends to an Exchange —
because the exchange decides where the message should go.

Technical Behavior

-> Opens a connection and a channel with RabbitMQ.
-> Publishes a message with:
- Exchange name
- Routing key
- Message body (payload)
- Properties (headers, persistence flag, priority, etc.)

every message sent by a producer reaches the exchange in binary form.
The client library transmits messages as binary (because AMQP is a binary protocol).
But it's the developer's responsibility to serialize structured data into bytes before sending.

Example

```
channel.publish('orders_exchange', 'order.created',
Buffer.from(JSON.stringify({
  orderId: 101,
  user: 'Aniket'
})), { persistent: true });
```

Responsibilities

Serialize data (e.g. JSON)
Define routing key
Ensure persistence and delivery confirmations

2. Exchange — the message router

The Exchange is the brain of RabbitMQ's routing system.
It decides how messages flow from producers to one or more queues.

When a message arrives, the exchange looks at its type and bindings to decide which queues should receive the message.

Main Exchange Types

Type	Routing Logic	Use Case
Direct	Exact routing key match	Logs, targeted routing
Fanout	Broadcasts to all queues bound	Notifications, broadcasts
Topic	Pattern-based routing	Multi-service routing (*.created, india.*)
Headers	Route based on headers	Advanced filtering

Responsibilities

Receive messages from producers

Match routing keys against bindings

Route copies of messages into the correct queues

Queue — the message buffer

A Queue is a storage unit that holds messages until they're consumed.
Think of it as a FIFO (first-in, first-out) message buffer.

Messages stay in the queue until:

-> A consumer retrieves and acknowledges them
-> The message expires (TTL)
-> The queue is deleted

Key Properties

Property	Description
Durable	Survives broker restarts
Exclusive	Used only by one connection
Auto-delete	Deleted when no consumers remain
TTL	Message expiry time
DLX	For failed messages
Lazy mode	Stores on disk to save RAM

Responsibilities

-> Hold messages safely until processed
-> Manage message order and delivery
-> Support reliability (ack, persistence, DLQ)
-> Support scaling (multiple consumers)

Queue Declaration

```
await channel.assertQueue(
'orderQueue' , {
  durable: true,
  exclusive: false,
  autoDelete: false,
  arguments: {}
});
```

Queue Types (Special Behavior Queues)

RabbitMQ doesn't have "named types" like "direct queue" or "topic queue" — the exchange defines routing,
but queue types here mean specialized behavior queues based on their arguments.

Let's explore the most important ones

1. Standard Queue (Default)

The most common queue type.

FIFO (First-In-First-Out) order is preserved (except when using priority).

Messages stored in memory or disk (depending on mode).

Can have multiple consumers (for load balancing).

Use case: General background jobs, microservices messaging.

2. Priority Queue

Messages can have a priority level (integer).

Higher-priority messages are delivered first, even if others arrived earlier.

Setup:

```
await channel.assertQueue('priorityQueue', {
  arguments: { 'x-max-priority': 10 }
});
```

Producer:

```
channel.sendToQueue('priorityQueue',
Buffer.from('Urgent task'),
{ priority: 8 });
```

Use case: Payment alerts, urgent notifications, critical updates.

3. Dead Letter Queue (DLQ)

Stores failed messages that couldn't be processed successfully.

Typically connected via a Dead Letter Exchange (DLX).

```
await channel.assertQueue('mainQueue', {
  arguments: {
    'x-dead-letter-exchange': 'dlx_exchange',
    'x-dead-letter-routing-key': 'dead.msg'
  }
});

await channel.assertQueue('dlqQueue');
await channel.bindQueue('dlqQueue', 'dlx_exchange',
'dead.msg');
```

Use case: Error handling, retries, message auditing.

4. Lazy Queue

Stores messages on disk immediately instead of memory.

Reduces RAM usage when message volume is high.

Slightly slower, but very memory efficient.

```
await channel.assertQueue('lazyQueue', {
  arguments: { 'x-queue-mode': 'lazy' }
});
```

Use case: Large queues, background workloads, high backlog systems.

5. TTL Queue (Time-To-Live)

Messages automatically expire after a certain time.

```
await channel.assertQueue('ttlQueue', {
  arguments: { 'x-message-ttl': 10000 } // 10 seconds
});
```

Use case: Temporary notifications, caching, expiring tasks.

6. Single Active Consumer Queue

Only one consumer is active at a time.
If it disconnects, another takes over.
Ensures message order and sequential processing.

```
await channel.assertQueue('sacQueue', {
  arguments: { 'x-single-active-consumer': true }
});
```

7. Stream Queue (Newer Feature)

Hybrid between queue and Kafka topic.

Messages are retained for a configurable time (like Kafka).

Consumers can replay messages by offset.
Use case: Event streaming, audit logging, replayable systems.

Binding — the connection between exchange and queue

A Binding is the rule that connects an exchange to a queue.
It defines which messages go into which queue, based on routing keys or headers.

Without bindings, exchanges can't route anything — they'd just drop messages.

Example

```
channel.bindQueue('emailQueue', 'order_exchange', 'order.email');
```

Whenever a producer sends a message to exchange order_exchange with routing key order.email,

It will be routed to emailQueue.

You can have:

One exchange bound to many queues
One queue bound to many exchanges
Multiple bindings for flexible routing

Example:

```
Exchange: order_events
 ├─ (key: order.created) → Queue: inventory
 └─ (key: order.created) → Queue: email
```

Consumer

A Consumer is an application or service that reads messages from a queue and processes them.

It connects to RabbitMQ, subscribes to a queue, and waits for new messages to arrive.

Consumers can:

Auto-acknowledge messages (less reliable)
Manually acknowledge after processing (recommended)
NACK messages to requeue or send to DLQ

channel.ack(msg)	Remove message from queue permanently
channel.nack(msg, false, true)	Requeue message for retry
channel.nack(msg, false, false)	Drop message (or send to DLQ)
(Consumer crash before ACK)	RabbitMQ requeues automatically

Responsibilities

Retrieve and process messages
Send ACK/NACK
Handle retries or DLQ
Scale horizontally for parallel processing

Example

```
channel.consume('orderQueue', (msg) => {
  const order = JSON.parse(msg.content.toString());
  console.log("Processing order:", order.orderId);
  // simulate processing
  channel.ack(msg);
});
```

Flow:

RabbitMQ delivers message → Consumer

Consumer processes it

Consumer sends ACK → RabbitMQ removes message

Bonus Internal Components (Under the Hood)

RabbitMQ has some hidden but critical internal components that make these 5 core parts work efficiently:

Channel

Lightweight virtual connection inside a TCP connection.

You can have multiple channels per connection.

Used for publishing, consuming, and managing queues.

Why?
Opening/closing TCP connections is expensive;
channels are cheap.

Example:
One app → one TCP connection →
multiple channels for producers/consumers.

Connection

A real TCP connection between your app and RabbitMQ broker.

Established using AMQP protocol (amqp://localhost:5672).

RabbitMQ allows multiple concurrent connections
from different apps or services.

Virtual Host (vHost)

Logical namespace for isolation.

Think of it like "databases" inside a DBMS.

Each vHost has its own exchanges, queues, and permissions.

Default vHost: /

Useful for multi-tenant systems or environment separation (/dev, /prod).

PRODUCER
↓
[Connection + Channel]
↓
EXCHANGE
↓
(Bindings)
↓
QUEUE
↓
CONSUMER

Message (Payload + Metadata)

```
{
  "body": "{ \"orderId\": 123 }",
  "properties": {
    "contentType": "application/json",
    "deliveryMode": 2,
    "headers": { "region": "india" }
  }
}
```

Direct Exchange

Routes messages to queues whose binding key matches the routing key exactly.

Example:

Message with key "error" → goes to queue bound with "error"

Message with key "info" → goes to "info" queue

DEMO

Fanout Exchange

Broadcasts messages to all queues bound to the exchange (ignores routing key).
Perfect for notifications or real-time broadcasts.

DEMO

Topic Exchange

Routes messages based on pattern matching of the routing key.

Routing keys are dot-separated words:

*.created → matches order.created, user.created

user.# → matches all events starting with user.

order.*
not order.created.success

order.#
also order.created.success

DEMO

Headers Exchange

Routes based on message headers, not routing key.
Supports x-match: all or x-match: any.

DEMO

Priority Queue

Priority in RabbitMQ only affects the order of messages that are already in the queue at the same time.

Normally, RabbitMQ delivers messages in the order they arrive (FIFO).

But sometimes, you want certain messages to jump ahead in the queue — for example:

"VIP" customer orders

"High urgency" background jobs

"Error" events over "Info" logs

DEMO

arguments: { "x-max-priority": 10 }

Priority values are integers from 0 (lowest) to max (highest).

channel.publish('exchange', 'key', Buffer.from(msg), { priority: 8 });

RabbitMQ stores higher-priority messages ahead of lower ones.

If two messages have the same priority → normal FIFO order applies.

AMQP (Advanced Message Queuing Protocol)

AMQP stands for Advanced Message Queuing Protocol —
It's the standard communication protocol used by RabbitMQ (and other brokers like Qpid, ActiveMQ).

Just like HTTP defines how browsers talk to servers,
AMQP defines how message producers, brokers, and consumers communicate.

RabbitMQ is a message broker implementation of this AMQP

AMQP Defines:

Layer	Purpose
Connection Layer	How clients connect/authenticate (TCP)
Session Layer (Channels)	How virtual channels are managed
Messaging Layer	How messages are routed, acknowledged, etc.

So AMQP is not just about "sending messages" — it defines exactly how:

Messages are published (with routing key, exchange)
Brokers store and forward messages
Consumers ACK/NACK messages
Reliability, persistence, and confirmations work

Example AMQP Command Flow

channel.publish("orders_exchange", "order.created", Buffer.from("data"));

Under the hood, RabbitMQ translates it into AMQP frames like:

Basic.Publish
→ Exchange: "orders_exchange"
→ Routing key: "order.created"
→ Properties: delivery_mode=2
→ Body: "data"

And when the consumer receives it:

Basic.Deliver
→ Queue: "orders"
→ Consumer Tag: 1234
→ Body: "data"

All this communication happens as binary AMQP frames over the TCP connection

Why AMQP Exists

Before AMQP, every message broker had its own proprietary protocol.
AMQP provides:

Interoperability — any AMQP-compliant client (Python, JS, Go, Java) can talk to any AMQP broker.

Reliability — standard support for acknowledgments, delivery guarantees, transactions, etc.

Security — authentication, permissions, virtual hosts.

Application Code (Producer / Consumer)
↓
Client Library (amqplib, pika, etc.)
↓
AMQP Protocol (frames, methods, acks)
↓
TCP Connection
↓
RabbitMQ Broker

SPOF in RabbitMQ

a standalone RabbitMQ node = SPOF

a single RabbitMQ node is a SPOF by default.
But you can make it highly available using clustering, mirrored queues, quorum queues or federation/shovel setups

Performance and Scalability

RabbitMQ is fast, but not as fast as Kafka.
Typical throughput: ~50K–100K msgs/sec/node (depends on hardware).
For millions msgs/sec → use Kafka

Ways to scale:

Horizontal Scaling: Add more nodes (RabbitMQ cluster)

Sharding: Divide queues among nodes

Consumers Scaling: Add more consumers for load balancing